



UNIVERSITY OF TORONTO  
SCHOOL OF CONTINUING STUDIES

# Doodle Recognition using Convolutional Neural Networks

SCS\_3253\_024 Machine Learning



# Group 8

**Nitin Prakash Panicker**



**Syed Reza**





# Hello!

*We want to explore the  
Quick-Draw dataset.*

Quick-draw is a 'neural network game, trying to guess what the user is drawing.



## How would you like to explore our project?



or, continue with  
this presentation...



First 2 inks will take you to external websites.



1.0

# Business Objectives

Why are we interested in this project?

Discovering this doodle dataset will  
extend opportunities to explore  
implementations in real life projects  
like, **handwritten** text recognition,  
touch based complex interactions, etc.



“

Knowledge from this project work will help understand challenges for human computer **interaction in touch-based applications.**

“





2.0

# The Dataset

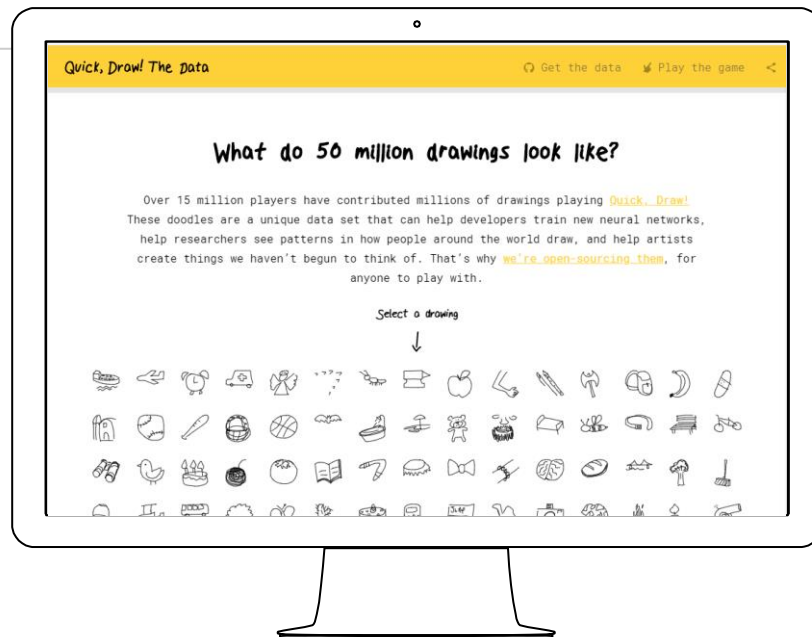
What do 50 million drawings look like?





## Quick, Draw! The Data

Millions of players have contributed huge doodle drawings playing this game! We decided to use these doodles to train our convolutional neural network.





## Introduction to the Dataset

- 50 million drawings
- Classified into 345 classes
- 20 classes of drawings with first 10,000 image samples taken for data analysis and wrangling
- For CNN data modelling, we have used all 345 classes, but first 10,000 images



## Where to **get** the data?

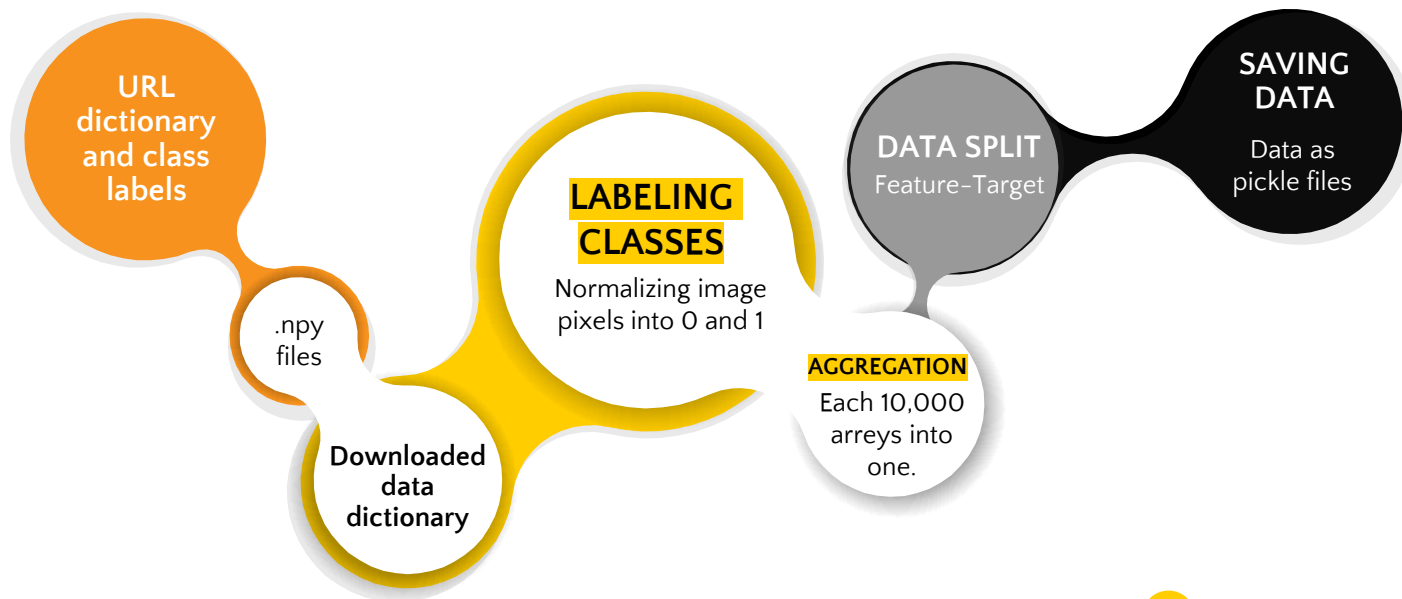
- Entire dataset is available in **GitHub**:  
<https://github.com/googlecreativelab/quickdraw-dataset>
- We are working with preprocessed dataset provided by Google Creative Lab.
- All the drawings have been rendered into a 28x28 px grayscale bitmap in numpy format.



## How data cleansing and wrangling were done?



## Data **Cleansing** Process



Link to  **GitHub source.**



## Data Wrangling

- Created a dictionary for URL and class labels to download the entire dataset in .npy files.
- Then the dataset was downloaded into a dictionary.
- Normalized the image pixels between 0 and 1 and added class labels.



## Data Wrangling

- Aggregated the first 10,000 arrays for every doodle into one array of doodles.
- Split the data into features and target labels.
- Save the train and test data as serialized data using pickle files.



Colab couldn't handle gigantic dataset!



Encountered a 'Memory Error' while trying to create a dictionary object using the downloaded data in Colab.



“



## Memory Error

In [130]: `%%time`

```
classes_dict = {}  
for key, value in URL_DATA.items():  
    response = requests.get(value)  
    classes_dict[key] = np.load(BytesIO(response.content))
```

-----  
**MemoryError**

Traceback (most recent call last)

<timed exec> in <module>

~\AppData\Local\Continuum\anaconda3\envs\nitin\lib\site-packages\requests\api.py in get(url, params, \*\*kwargs)

74

75 kwargs.setdefault('allow\_redirects', True)

---> 76 return request('get', url, params=params, \*\*kwargs)

77

78

Running the code in Colab crashed the application and produced a **RAM error** as well. Then solved by looping through every .npy file and .png image data for every .npy file to set up data for 345 samples to create our CNN model.



“



## Generating PNG by **Looping** all .npy

```
%%time

# Looping through every npy file and generating png image data for every npy file.
for filename in os.listdir(data_file_path):
    if filename.endswith('.npy'):
        data = np.load(data_file_path+"/"+filename)
        print(filename)
        data = data[0,:]

    if not os.path.exists(train_path+os.path.splitext(filename)[0]):
        try:
            os.makedirs(train_path+os.path.splitext(filename)[0])
        except OSError as exc:
            if exc.errno != errno.EEXIST:
                raise
    print(train_path+os.path.splitext(filename)[0])

    if not os.path.exists(test_path+os.path.splitext(filename)[0]):
        try:
            os.makedirs(test_path+os.path.splitext(filename)[0])
        except OSError as exc:
            if exc.errno != errno.EEXIST:
                raise
    print(test_path+os.path.splitext(filename)[0])

    for i in range(0,5):
        x=np.reshape(data[i],(28,28))
        img = Image.fromarray(x)
        img = img.convert('L')
        if i<4:
            img.save(train_path+os.path.splitext(filename)[0+"/"+os.path.splitext(filename)[0]+str(i)+".png")
        else:
            img.save(test_path+os.path.splitext(filename)[0+"/"+os.path.splitext(filename)[0]+str(i-4)+".png")
```



## Exploring the hidden layer



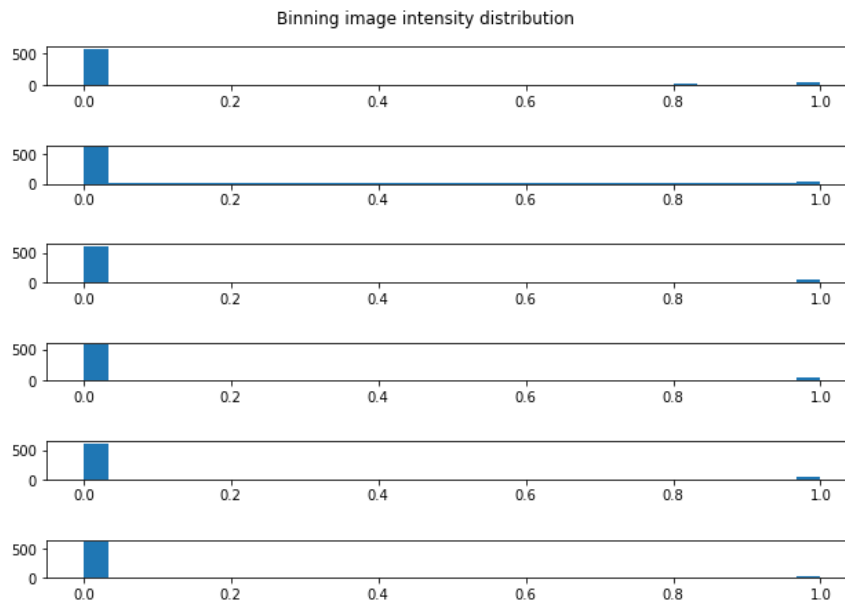
## Exploring Test Data

```
In [44]: ▶ # Test set  
         unique, counts = np.unique(y_test, return_counts=True)  
         dict(zip(unique, counts))
```

```
Out[44]: {0.0: 2027,  
          1.0: 2062,  
          2.0: 2081,  
          3.0: 2051,  
          4.0: 2000,  
          5.0: 2075,  
          6.0: 1974,  
          7.0: 2012,  
          8.0: 1984,  
          9.0: 1946,  
          10.0: 2034,  
          11.0: 1966,  
          12.0: 1967,  
          13.0: 1997,  
          14.0: 1921,  
          15.0: 1993,  
          16.0: 1984,  
          17.0: 1901,  
          18.0: 2010,  
          19.0: 2015}
```



# Plotting Intensity Histogram





## Building the CNN model





## CNN Model Architecture for 345 Classes

```
In [110]: model = Sequential([Convolution2D(filters = 32,
      kernel_size = (3,3),
      activation = "relu",
      input_shape = img_size),
      Convolution2D(filters = 32,
      kernel_size = (3,3),
      activation = "relu"),
      MaxPooling2D(pool_size=(2,2)),
      Dropout(.20),
      Convolution2D(filters = 64,
      kernel_size = (3,3),
      activation = "relu"),
      Convolution2D(filters = 64,
      kernel_size = (3,3),
      activation = "relu"),
      MaxPooling2D(pool_size=(2,2)),
      Dropout(.1),
      Flatten(),
      Dense(512, activation="relu"),
      Dense(n_classes, activation="softmax")
    ])
```



## Model Summary

```
model.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
conv2d_21 (Conv2D)	(None, 26, 26, 32)	320
conv2d_22 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_11 (MaxPooling)	(None, 12, 12, 32)	0
dropout_11 (Dropout)	(None, 12, 12, 32)	0
conv2d_23 (Conv2D)	(None, 10, 10, 64)	18496
conv2d_24 (Conv2D)	(None, 8, 8, 64)	36928
max_pooling2d_12 (MaxPooling)	(None, 4, 4, 64)	0
dropout_12 (Dropout)	(None, 4, 4, 64)	0
flatten_6 (Flatten)	(None, 1024)	0
dense_15 (Dense)	(None, 512)	524800
dense_16 (Dense)	(None, 345)	176985
Total params: 766,777		
Trainable params: 766,777		
Non-trainable params: 0		



## Model Training

```
In [99]: %%time

training_set = train_datagen.flow_from_directory(
    train_path,
    target_size=(28, 28),
    color_mode="grayscale",
    batch_size=256,
    class_mode='categorical')
```

Found 2760000 images belonging to 345 classes.  
Wall time: 1min 39s

```
In [100]: %%time

test_set = test_datagen.flow_from_directory(
    test_path,
    target_size=(28, 28),
    color_mode="grayscale",
    batch_size=256,
    class_mode='categorical')
```

Found 690000 images belonging to 345 classes.  
Wall time: 28.6 s



## Model Fitting

```
history = model.fit(  
    training_set,  
    steps_per_epoch=training_set.n//training_set.batch_size,  
    epochs=15,  
    validation_data=test_set,  
    validation_steps=test_set.n//test_set.batch_size,  
    callbacks=[save_cb, TensorBoard(log_dir="QuickDraw")])
```

```
#history = model.fit(train_x, train_y, validation_data=(test_x, test_y), epochs=5, batch_size=256,  
#                    callbacks=[TensorBoard(log_dir="QuickDraw"), save_cb])
```



## Do You See the Wall Time?

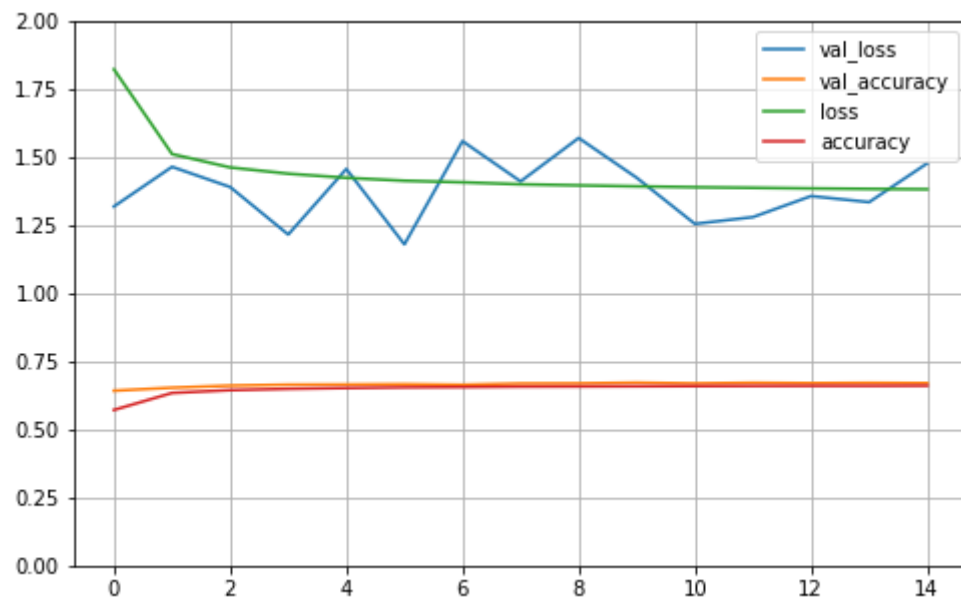
```
10781/10781 [=====] - 5197s 482ms/step - loss: 1.3964 - accuracy: 0.6578 - val_lo  
ss: 1.5704 - val_accuracy: 0.6688  
Epoch 10/15  
10781/10781 [=====] - 5212s 483ms/step - loss: 1.3928 - accuracy: 0.6584 - val_lo  
ss: 1.4241 - val_accuracy: 0.6715  
Epoch 11/15  
10781/10781 [=====] - 5153s 478ms/step - loss: 1.3897 - accuracy: 0.6591 - val_lo  
ss: 1.2549 - val_accuracy: 0.6686  
Epoch 12/15  
10781/10781 [=====] - 5066s 470ms/step - loss: 1.3874 - accuracy: 0.6597 - val_lo  
ss: 1.2798 - val_accuracy: 0.6711  
Epoch 13/15  
10781/10781 [=====] - 5166s 479ms/step - loss: 1.3851 - accuracy: 0.6601 - val_lo  
ss: 1.3570 - val_accuracy: 0.6697  
Epoch 14/15  
10781/10781 [=====] - 6328s 587ms/step - loss: 1.3834 - accuracy: 0.6601 - val_lo  
ss: 1.3347 - val_accuracy: 0.6707  
Epoch 15/15  
10781/10781 [=====] - 5627s 522ms/step - loss: 1.3821 - accuracy: 0.6606 - val_lo  
ss: 1.4774 - val_accuracy: 0.6698  
Wall time: 1d 50min 55s
```

*Model fitting was running little more than a day. The machine did not have enough hardware to contribute to the processing time.*

“



## Model Performance



Our CNN model did not perform well enough. We ended up with a validation accuracy of ~70%. Which isn't very attractive. Inclusion of all 345 classes might have been too ambitious for 10,000 sample images.



“





5.0

## Lessons Learned

What could have been better

Selecting *less than 345 classes* could have shown a better result with added opportunity of testing several models.

“

Touch based human computer interaction for **handicapped users** could be improved more with exploratory datasets from various world regions.

“



UNIVERSITY OF TORONTO  
SCHOOL OF CONTINUING STUDIES



# Thanks!

## Any **questions** ?

Please let us know in



Quercus