



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Doodle Recognition using Convolutional Neural Networks

SCS_3253_024 Machine Learning



Group 8

Nitin Prakash Panicker



Syed Reza





Hello!

*We want to explore the
Quick-Draw dataset.*

Quick-draw is a 'neural network game, trying to guess what the user is drawing.



This is what **Quick, Draw!** Look Like



✓ map



✓ dishwasher



✓ broccoli



✓ mermaid



✓ airplane



✗ belt



Clicking on this image will take you to the game. 4



How would you like to explore our project?



or, continue with
this presentation...



First 2 inks will take you to external websites.



Why are we interested in this project?

Discovering this doodle dataset will
extend opportunities to explore
implementations in real life projects
like, **handwritten** text recognition,
touch based complex interactions, etc.



“

Knowledge from this project work will help understand challenges for human computer **interaction in touch-based applications.**

“





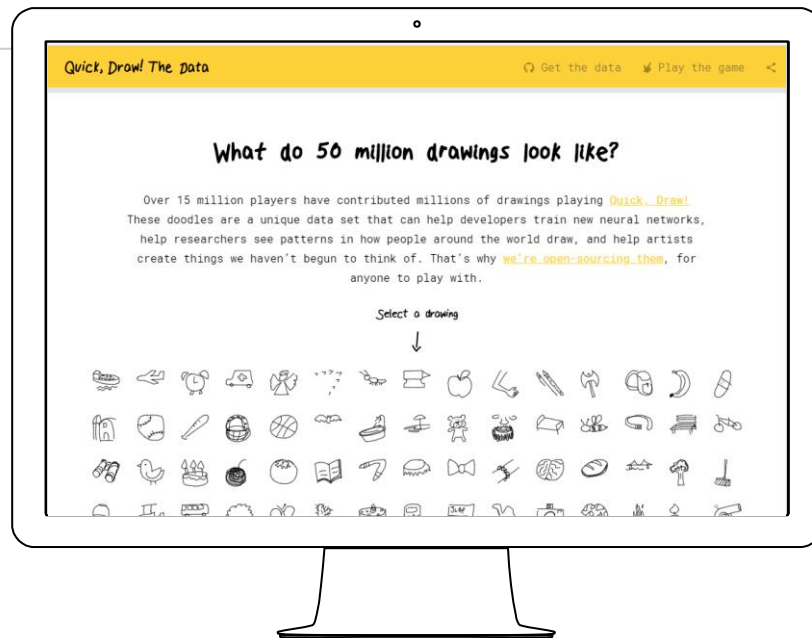
2.0

The Dataset

What do 50 million drawings look like?



Millions of players have contributed huge doodle drawings playing this game! We decided to use these doodles to train our convolutional neural network.





Introduction to the Dataset

- 50 million drawings
- Classified into 345 classes
- 20 classes of drawings with first 10,000 image samples taken for data analysis and wrangling
- For CNN data modelling, we have used all 345 classes, but first 10,000 images



Where to **get** the data?


- Entire dataset is available in **GitHub**:
<https://github.com/googlecreativelab/quickdraw-dataset>
- We are working with preprocessed dataset provided by Google Creative Lab.
- All the drawings have been rendered into a 28x28 px grayscale bitmap in numpy format.



How data cleansing and wrangling were done?



Data **Cleansing** Process



Link to  GitHub source.



Data Wrangling

- Created a dictionary for URL and class labels to download the entire dataset in .npy files.
- Then the dataset was downloaded into a dictionary.
- Normalized the image pixels between 0 and 1 and added class labels.



Data Wrangling

- Aggregated the first 10,000 arrays for every doodle into one array of doodles.
- Split the data into features and target labels.
- Save the train and test data as serialized data using pickle files.



Colab couldn't handle gigantic dataset!

Encountered a 'Memory Error' while trying to create a dictionary object using the downloaded data in Colab.



“



Memory Error

In [130]: `%%time`

```
classes_dict = {}  
for key, value in URL_DATA.items():  
    response = requests.get(value)  
    classes_dict[key] = np.load(BytesIO(response.content))
```

MemoryError

Traceback (most recent call last)

<timed exec> in <module>

~\AppData\Local\Continuum\anaconda3\envs\nitin\lib\site-packages\requests\api.py in get(url, params, **kwargs)

74

75 kwargs.setdefault('allow_redirects', True)

---> 76 return request('get', url, params=params, **kwargs)

77

78

Running the code in Colab crashed the application and produced a **RAM error** as well. Then solved by looping through every .npy file and .png image data for every .npy file to set up data for 345 samples to create our CNN model.



“



Generating PNG by **Looping** all .npy

```
%%time

# Looping through every npy file and generating png image data for every npy file.
for filename in os.listdir(data_file_path):
    if filename.endswith('.npy'):
        data = np.load(data_file_path+"/"+filename)
        print(filename)
        data = data[0,:]

    if not os.path.exists(train_path+os.path.splitext(filename)[0]):
        try:
            os.makedirs(train_path+os.path.splitext(filename)[0])
        except OSError as exc:
            if exc.errno != errno.EEXIST:
                raise
    print(train_path+os.path.splitext(filename)[0])

    if not os.path.exists(test_path+os.path.splitext(filename)[0]):
        try:
            os.makedirs(test_path+os.path.splitext(filename)[0])
        except OSError as exc:
            if exc.errno != errno.EEXIST:
                raise
    print(test_path+os.path.splitext(filename)[0])

    for i in range(0,5):
        x=np.reshape(data[i],(28,28))
        img = Image.fromarray(x)
        img = img.convert('L')
        if i<4:
            img.save(train_path+os.path.splitext(filename)[0]+"/"+os.path.splitext(filename)[0]+str(i)+".png")
        else:
            img.save(test_path+os.path.splitext(filename)[0]+"/"+os.path.splitext(filename)[0]+str(i-4)+".png")
```



3.0

Data Analysis

Exploring the hidden layer



Exploring Test Data

```
In [19]: classes = ['star', 'apple', 'fish', 'moon', 'ice cream', 'tent', 'candle', 'fork', 'crab',  
                  'cake', 'hammer', 'sun', 'diamond', 'table', 'light bulb', 'spider', 'face', 'eye',  
                  'crown', 'shorts']  
  
#for filename in os.listdir("datasets/data/test_set"):  
#    classes.append(os.path.splitext(filename)[0])
```



Exploring Test Data

```
In [44]: ▶ # Test set  
         unique, counts = np.unique(y_test, return_counts=True)  
         dict(zip(unique, counts))
```

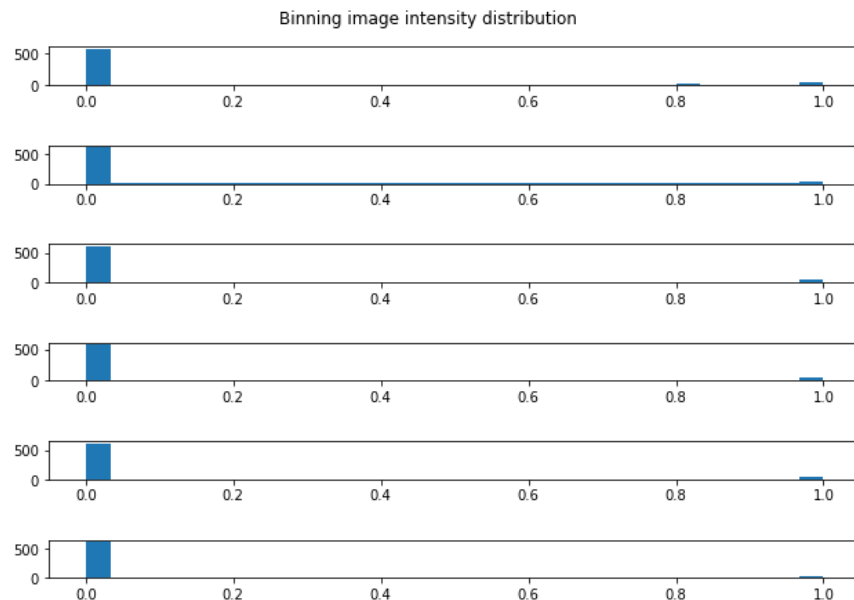

```
Out[44]: {0.0: 2027,  
          1.0: 2062,  
          2.0: 2081,  
          3.0: 2051,  
          4.0: 2000,  
          5.0: 2075,  
          6.0: 1974,  
          7.0: 2012,  
          8.0: 1984,  
          9.0: 1946,  
          10.0: 2034,  
          11.0: 1966,  
          12.0: 1967,  
          13.0: 1997,  
          14.0: 1921,  
          15.0: 1993,  
          16.0: 1984,  
          17.0: 1901,  
          18.0: 2010,  
          19.0: 2015}
```




Plotting Intensity Histogram

Image Intensity Histogram for each class

```
fig, axs = plt.subplots(20, 2, figsize=(15,20))
plt.suptitle('Binning image intensity distribution', x=.7,y=1)
for i, (key,value) in enumerate(classes_dict.items()):
    axs[i,0].set_title(key.upper())
    axs[i,0].imshow(value[0,:784].reshape((28,28)), cmap="gray_r")
    axs[i,1].hist(value[0,:784].reshape((28,28)).flatten(), bins=30,)
plt.tight_layout()
```





4.0

Data Model

Building the CNN model



CNN Model Architecture for 345 Classes

```
In [110]: model = Sequential([Convolution2D(filters = 32,
      kernel_size = (3,3),
      activation = "relu",
      input_shape = img_size),
      Convolution2D(filters = 32,
      kernel_size = (3,3),
      activation = "relu"),
      MaxPooling2D(pool_size=(2,2)),
      Dropout(.20),
      Convolution2D(filters = 64,
      kernel_size = (3,3),
      activation = "relu"),
      Convolution2D(filters = 64,
      kernel_size = (3,3),
      activation = "relu"),
      MaxPooling2D(pool_size=(2,2)),
      Dropout(.1),
      Flatten(),
      Dense(512, activation="relu"),
      Dense(n_classes, activation="softmax")
  ])
```



Model Summary

```
model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
conv2d_21 (Conv2D)	(None, 26, 26, 32)	320
conv2d_22 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_11 (MaxPooling)	(None, 12, 12, 32)	0
dropout_11 (Dropout)	(None, 12, 12, 32)	0
conv2d_23 (Conv2D)	(None, 10, 10, 64)	18496
conv2d_24 (Conv2D)	(None, 8, 8, 64)	36928
max_pooling2d_12 (MaxPooling)	(None, 4, 4, 64)	0
dropout_12 (Dropout)	(None, 4, 4, 64)	0
flatten_6 (Flatten)	(None, 1024)	0
dense_15 (Dense)	(None, 512)	524800
dense_16 (Dense)	(None, 345)	176985
Total params: 766,777		
Trainable params: 766,777		
Non-trainable params: 0		



Model Training

```
train_datagen2 = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    rescale=1./255,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    vertical_flip=False,  
    fill_mode='nearest')  
  
test_datagen2 = ImageDataGenerator(rescale=1./255)
```

```
In [99]: %%time  
  
training_set = train_datagen.flow_from_directory(  
    train_path,  
    target_size=(28, 28),  
    color_mode="grayscale",  
    batch_size=256,  
    class_mode='categorical')
```

Found 2760000 images belonging to 345 classes.
Wall time: 1min 39s

```
In [100]: %%time  
  
test_set = test_datagen.flow_from_directory(  
    test_path,  
    target_size=(28, 28),  
    color_mode="grayscale",  
    batch_size=256,  
    class_mode='categorical')
```

Found 690000 images belonging to 345 classes.
Wall time: 28.6 s



Model Optimizing

```
model.compile(optimizer = "nadam", loss="categorical_crossentropy", metrics=["accuracy"])
```

```
save_cb = ModelCheckpoint("conv_doodle_model.h5",  
                           save_best_only=True)
```



Model Fitting

```
history = model.fit(  
    training_set,  
    steps_per_epoch=training_set.n//training_set.batch_size,  
    epochs=15,  
    validation_data=test_set,  
    validation_steps=test_set.n//test_set.batch_size,  
    callbacks=[save_cb, TensorBoard(log_dir="QuickDraw")])
```



Do You See the Wall Time?

```
10781/10781 [=====] - 5197s 482ms/step - loss: 1.3964 - accuracy: 0.6578 - val_lo  
ss: 1.5704 - val_accuracy: 0.6688  
Epoch 10/15  
10781/10781 [=====] - 5212s 483ms/step - loss: 1.3928 - accuracy: 0.6584 - val_lo  
ss: 1.4241 - val_accuracy: 0.6715  
Epoch 11/15  
10781/10781 [=====] - 5153s 478ms/step - loss: 1.3897 - accuracy: 0.6591 - val_lo  
ss: 1.2549 - val_accuracy: 0.6686  
Epoch 12/15  
10781/10781 [=====] - 5066s 470ms/step - loss: 1.3874 - accuracy: 0.6597 - val_lo  
ss: 1.2798 - val_accuracy: 0.6711  
Epoch 13/15  
10781/10781 [=====] - 5166s 479ms/step - loss: 1.3851 - accuracy: 0.6601 - val_lo  
ss: 1.3570 - val_accuracy: 0.6697  
Epoch 14/15  
10781/10781 [=====] - 6328s 587ms/step - loss: 1.3834 - accuracy: 0.6601 - val_lo  
ss: 1.3347 - val_accuracy: 0.6707  
Epoch 15/15  
10781/10781 [=====] - 5627s 522ms/step - loss: 1.3821 - accuracy: 0.6606 - val_lo  
ss: 1.4774 - val_accuracy: 0.6698  
Wall time: 1d 50min 55s
```



Model fitting was running little more than **a day**. The machine did not have enough hardware to contribute to the processing time.



“



Model Performance



Our CNN model did not perform well enough. We ended up with a validation accuracy of ~70%. Which isn't very attractive. Inclusion of all 345 classes might have been too ambitious for 10,000 sample images.



“



5.0

Lessons Learned

What could have been better

Selecting *less than 345 classes* could have shown a better result with added opportunity of testing several models.

“

Touch based human computer interaction for **handicapped users** could be improved more with exploratory datasets from various world regions.



“



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES



Thanks!

Any **questions** ?

Please let us know in  Quercus